



Stony Brook University

# Adaptive Transaction Management in Permissioned Blockchains

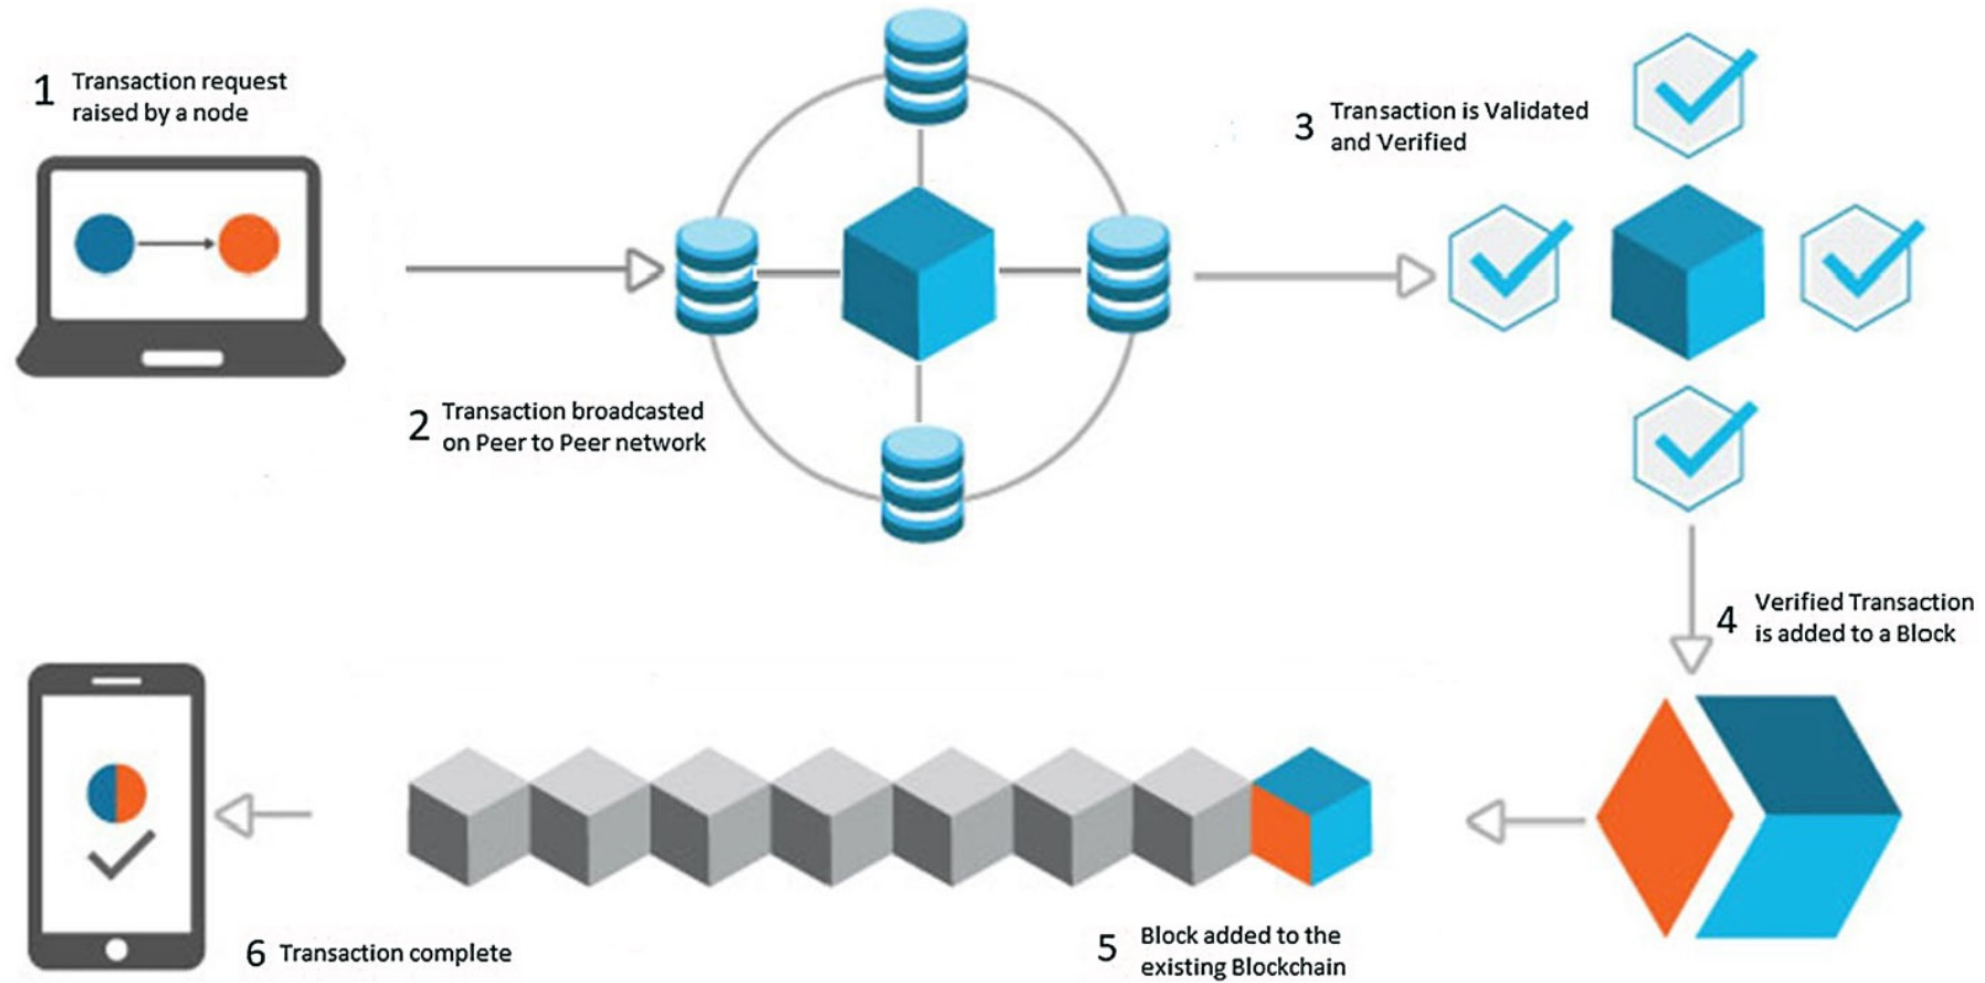
Sai Praneeth Reddy, Siddhartha Malladi, Madhu Lakkoju

# Contents

1. Introduction
2. Problem
3. Existing Solution
4. Proposed Solution
5. System Design
6. Learning Agent and Algorithms
7. Implementation
8. Results
9. Conclusion
10. References



# Blockchain



# Problem

- **Blockchain as a Service offering has a vast variety of use cases that lead to different workloads.**
- Workload Characteristics:
  - varying read/write ratios
  - skewness of popular key
  - compute intensity
- Not adaptable to Different workloads
- Lead to poor performance and inefficient resource utilization.
- If we use some heuristics, if hardware configurations change the heuristics might fail.



# Workloads:

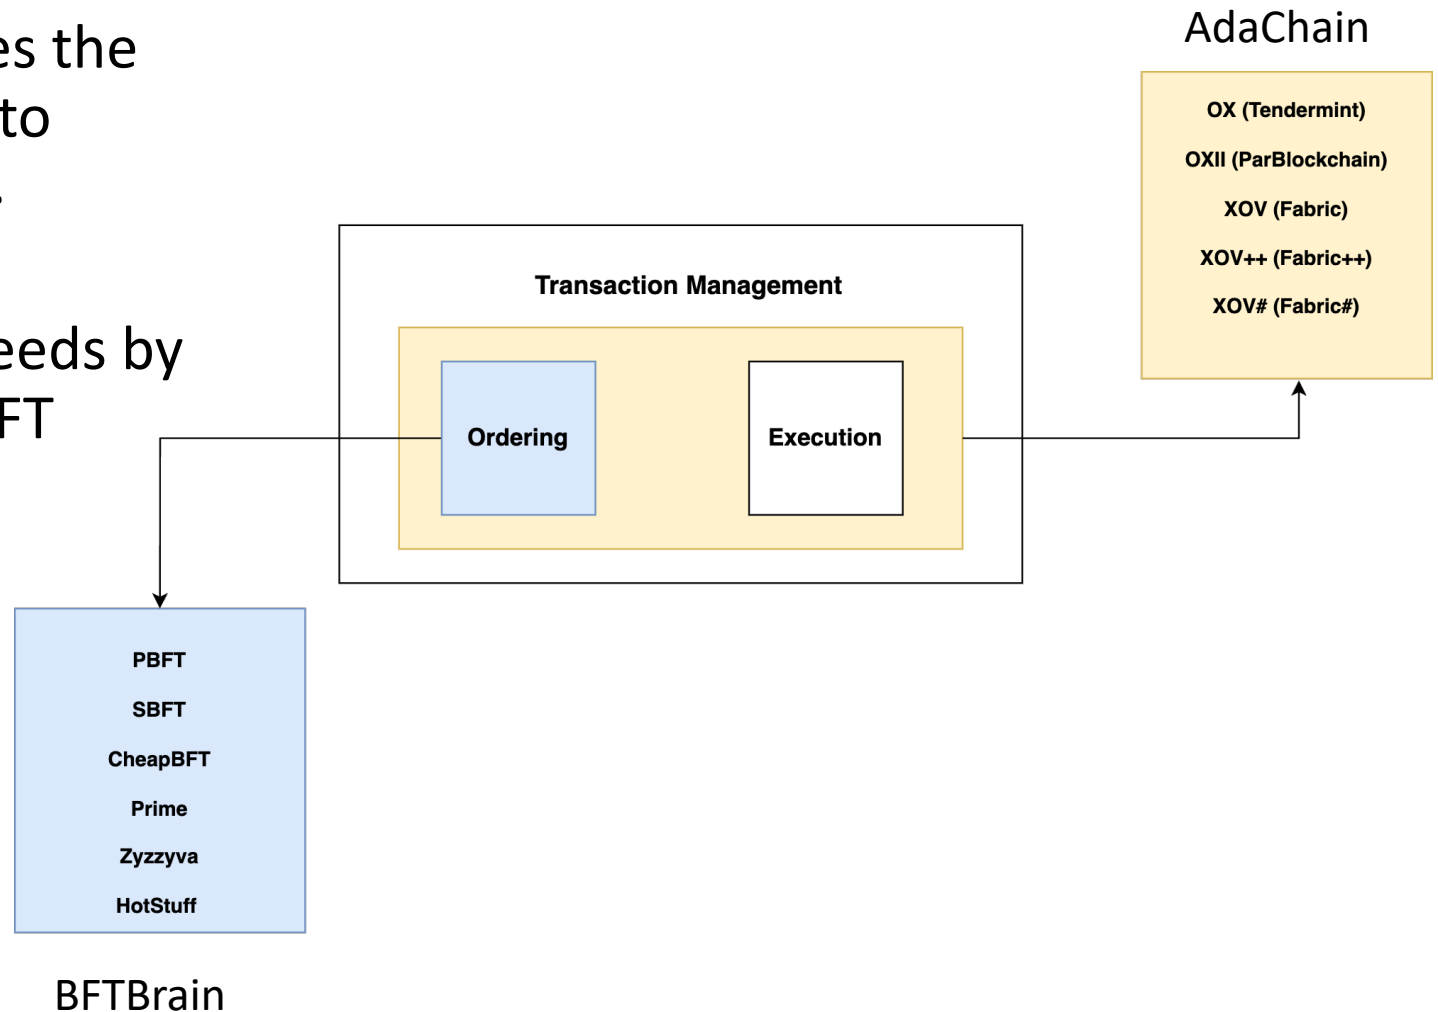
Workload	f	Write Ratio	Contention Level	Load	Compute Intensity	Req. size	Proposal slowness	Absent count
A	1	Low	High	High	High	1 KB	0ms	0
B	1	Mid	High	Low	Mid	1 KB	100 ms	0
C	4	Mid	Mid	High	Low	4 KB	20 ms	4
D	4	High	Low	High	High	100 KB	0 ms	0

Different Workloads demand different requirements like compute, parallel execution, and block size  
So, no one could identify the optimal configuration.




# Existing Solutions

- **AdaChain:** Adaptively chooses the best blockchain architecture to produce optimal throughput.
- **BFT Brain:** Adapts to system conditions and application needs by switching between a set of BFT protocols in real time.



# Proposed Solution: Fully Adaptive Transaction Management System

- This automatically **switches** between configurations (Architecture and BFT Protocol) to optimize performance online, compensating for changes in **workload** and **network conditions** to **maximize throughput**.
- Liveness, Safety.
- A fully decentralized machine-learning approach.
- Hardware change resistant
- Give best performance



# Configuration Landscape


- **Architecture Pool:**

- OX (Tendermint)
- OXII (ParBlockchain)
- XOY (Fabric)
- XOY++ (Fabric++)
- XOY# (Fabric#)


**Performance Params:** Block Size, Early Execution, Dependency Graph, Early Abort, Parallel Execution

- **Protocol Pool:**

- PBFT
- SBFT
- CheapBFT
- Prime
- Zyzzyva
- HotStuff-2




SBFT



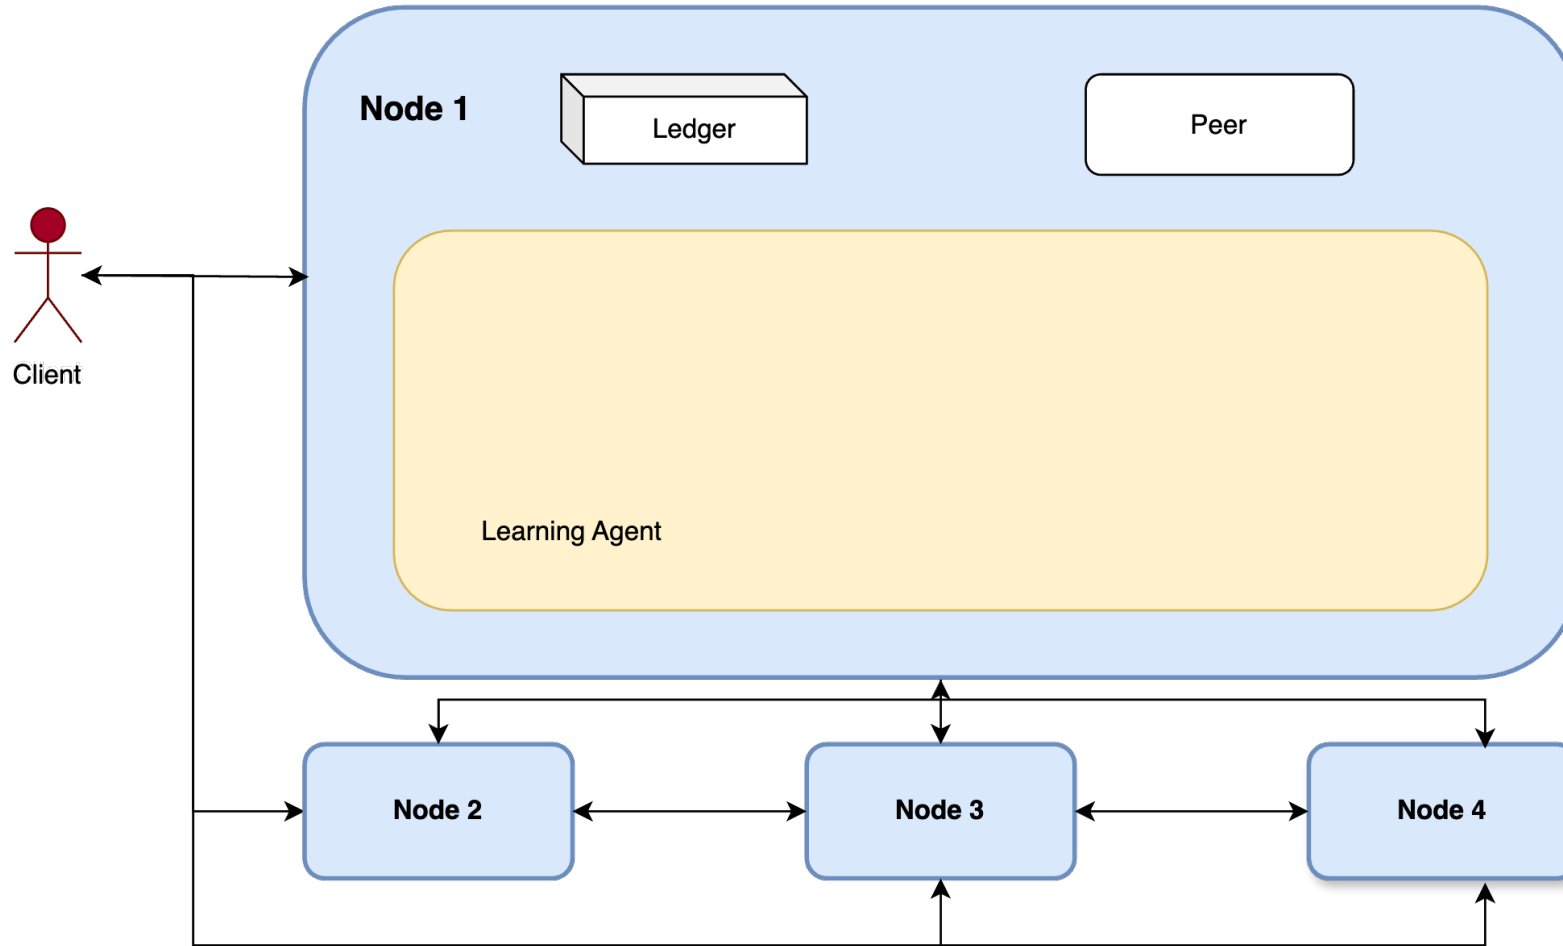
CheapBFT



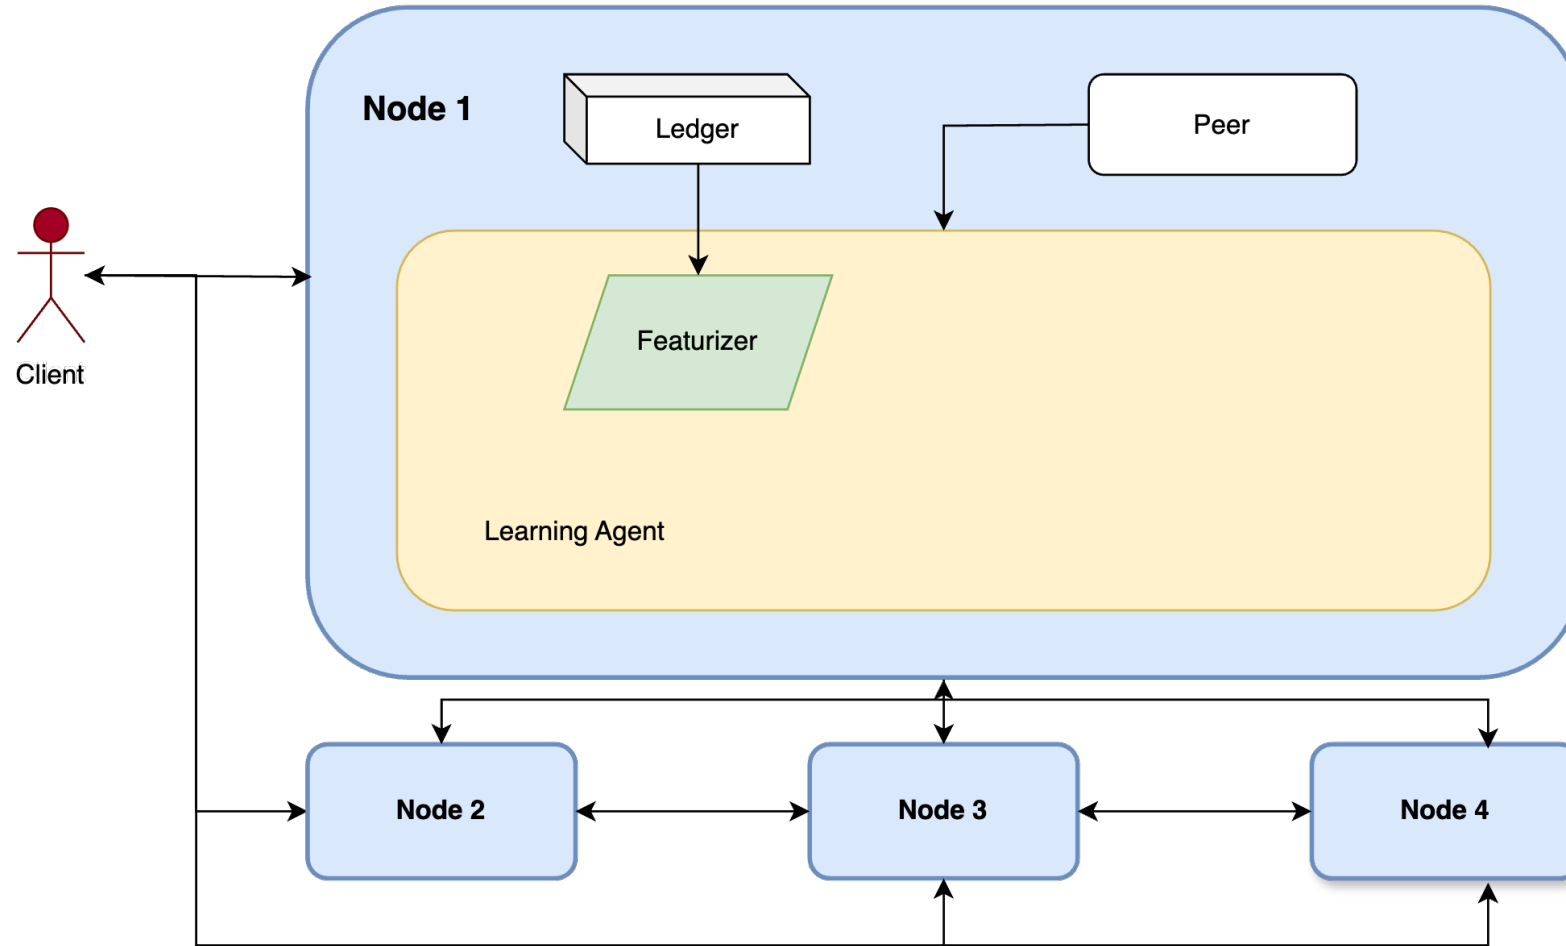
# System Design



# System Design

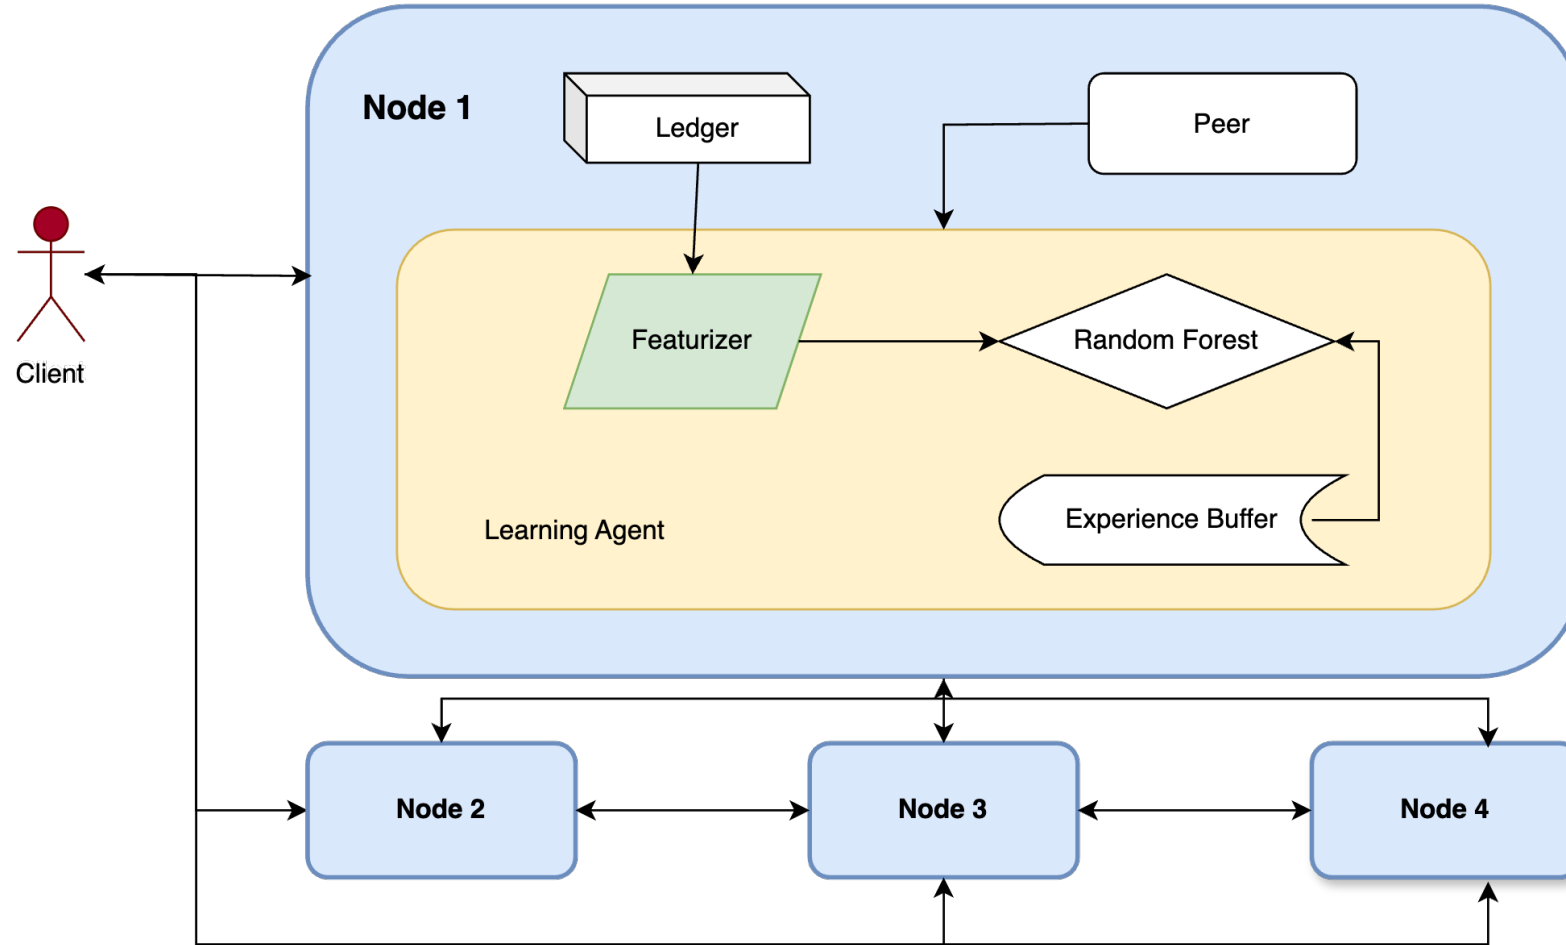


# System Design



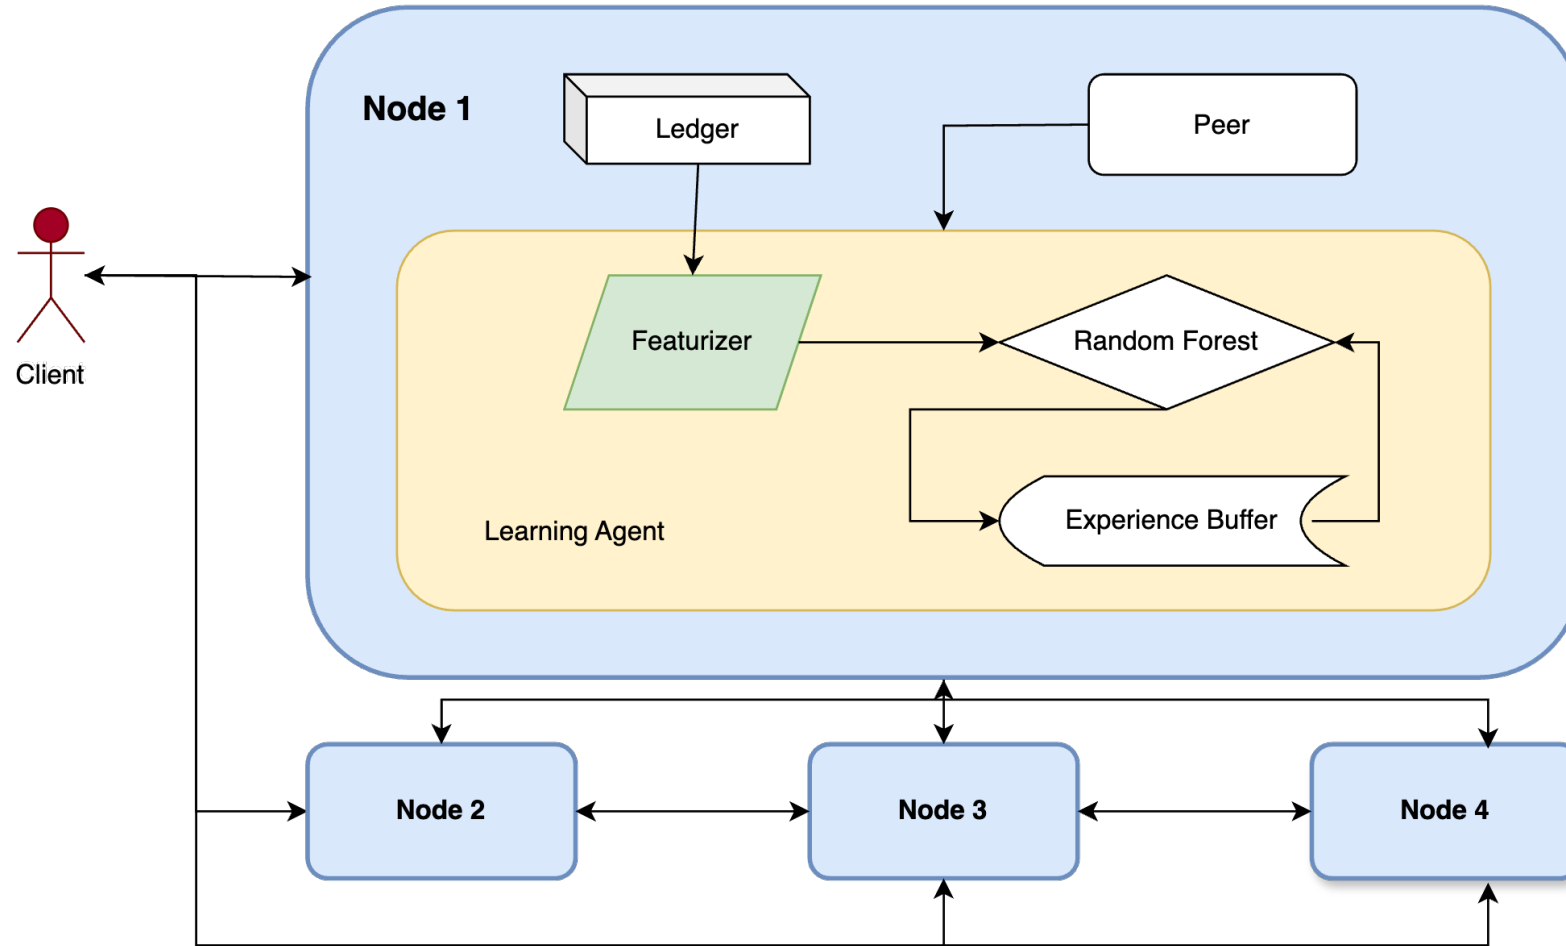
1. Notifying the learning agent

# System Design



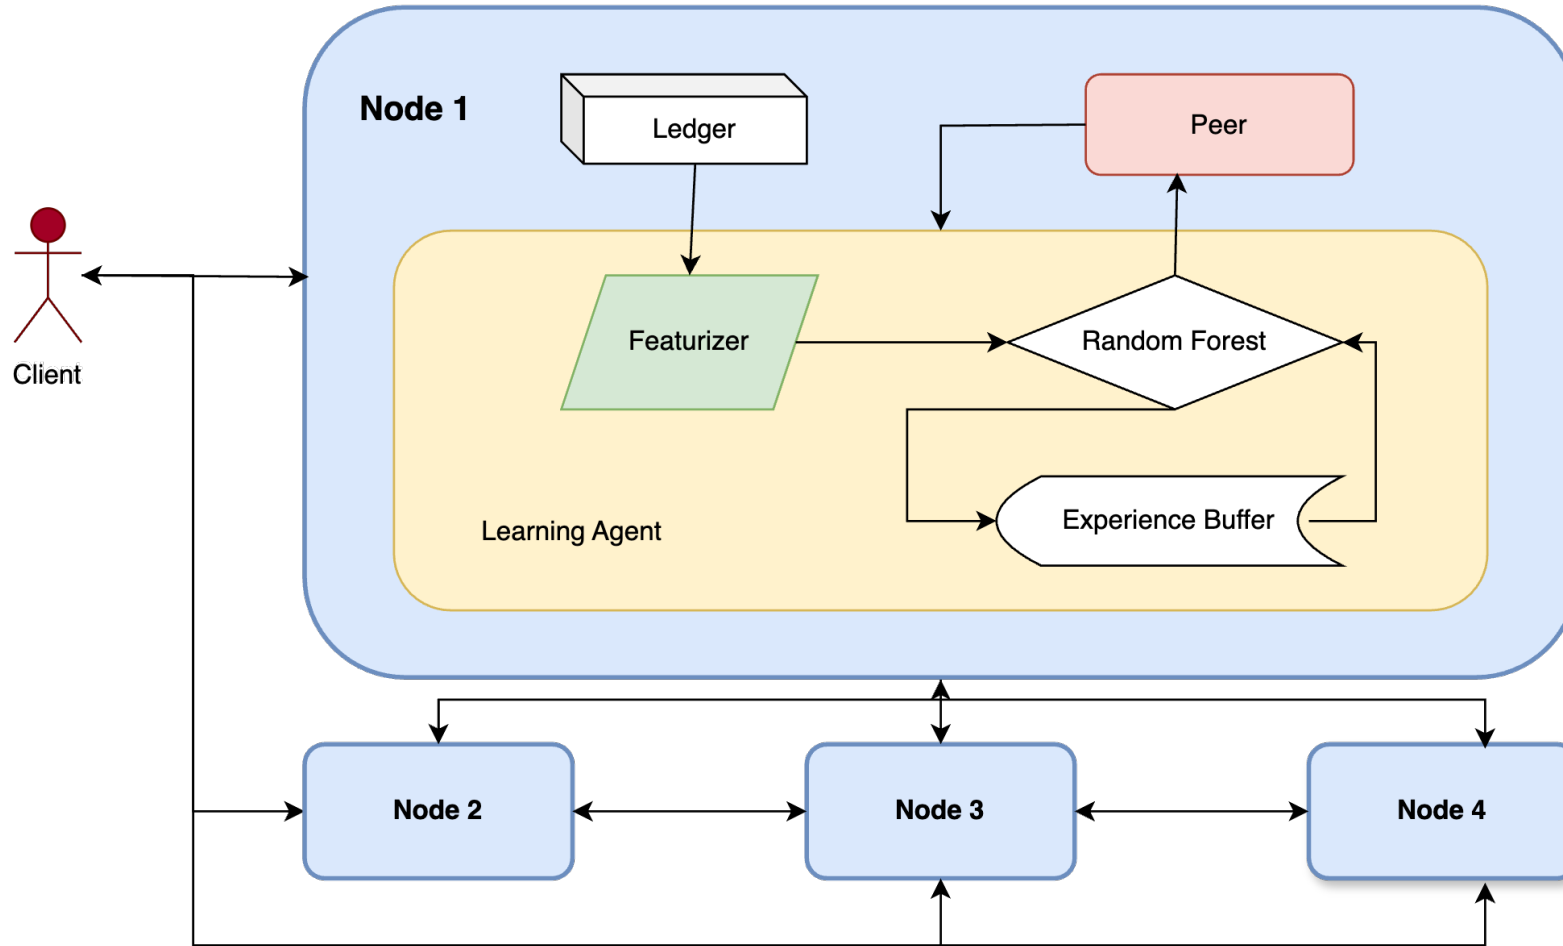
1. Notifying the learning agent
2. Featurization

# System Design



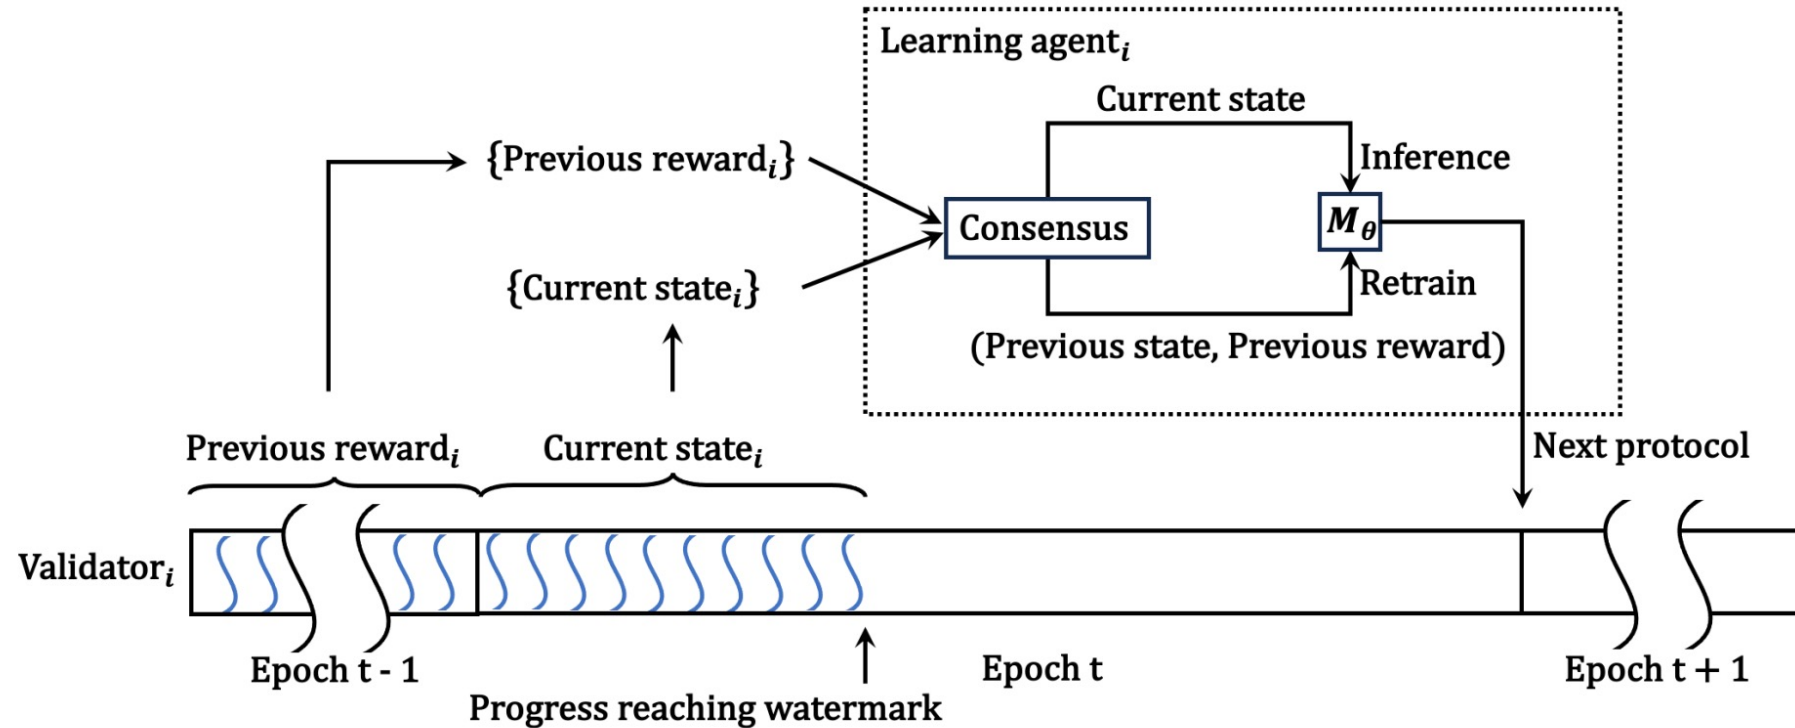
1. Notifying the learning agent
2. Featurization
3. Exchanging performance metrics
4. Estimating the performance of each architecture

# System Design




1. Notifying the learning agent
2. Featurization
3. Exchanging performance metrics
4. Estimating the performance of each architecture
5. Building experience buffer
6. Retraining

# Episodes: Watermarks and Actions



# Reinforcement Learning Approach:

- Learns from its **mistakes** and self-corrects through trials.
- To develop a good heuristic, an expert needs to **exhaustively experiment, not needing upfront data.**
- **Why not a supervised learning model?** - Supervised learning that assumes **training data is complete** and requires a separate data collection process before **deployment**
- If hardware changes full re-training is required





# Learning Agent and Algorithms

- We leverage reinforcement learning.
- Contextual multi-armed bandit problem – Minimize regret
  - **Context:** Workload
  - **Arms:** Configurations
  - **Reward:** Effective throughput
- **Problem Formulation:**  $r_n = (\text{optimal configuration} - \text{selected configuration})$
- **Thompson Sampling** - Balances Exploitation & Exploration
- **Predictive model** - Random Forest (with Thompson Sampling)
  - **INPUT:** Workload, configurations
  - **OUTPUT:** performance
- **Learning Coordination:** Learning Agent choose the best config, consensus between the learning agent is done.



# State Space and Action Space

- All aborted or invalidated transactions are still written to the ledger with a **validity flag**
- **State Space:** Helps in featurizing the current state, measurement in between 1 block
  - Write a ratio: Ratio of write transactions
  - Hot-key ratio
  - Transaction Arrival Rate
  - Execution Delay
- **Action space:**
  - Consists of a set of blockchain architectures and BFT Consensus protocols to choose from.

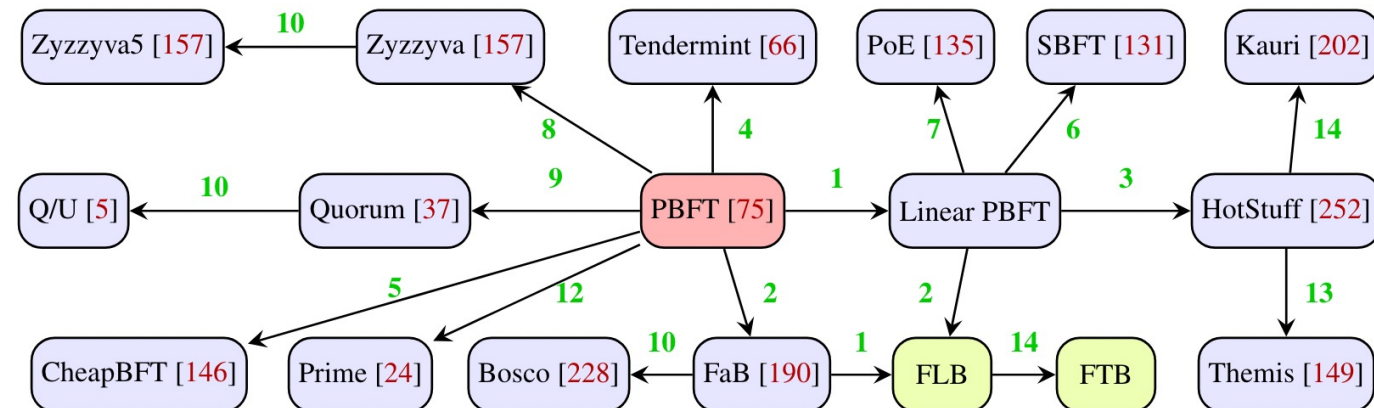

# Switching Configurations – Safety and Liveness

- **Normal Path:** Leader calculates median throughput
- **Slow Path:**
  - Bad configs, stalled commits, Timeouts happen
  - Stop block formation and create a consensus
  - Helpful in breaking bad configs, Consistency
- **Liveness:** Switch configs in a distributed fashion, without stalling the system.



# Bedrock

- Outstanding Paper NSDI 2024
- Plug and Play system
- A platform for :
  - BFT protocol **analysis**
  - BFT Protocols **implementation**
  - BFT Protocols **experimentation**.
- Capture the trade-offs between different design space dimensions.



# Implementation

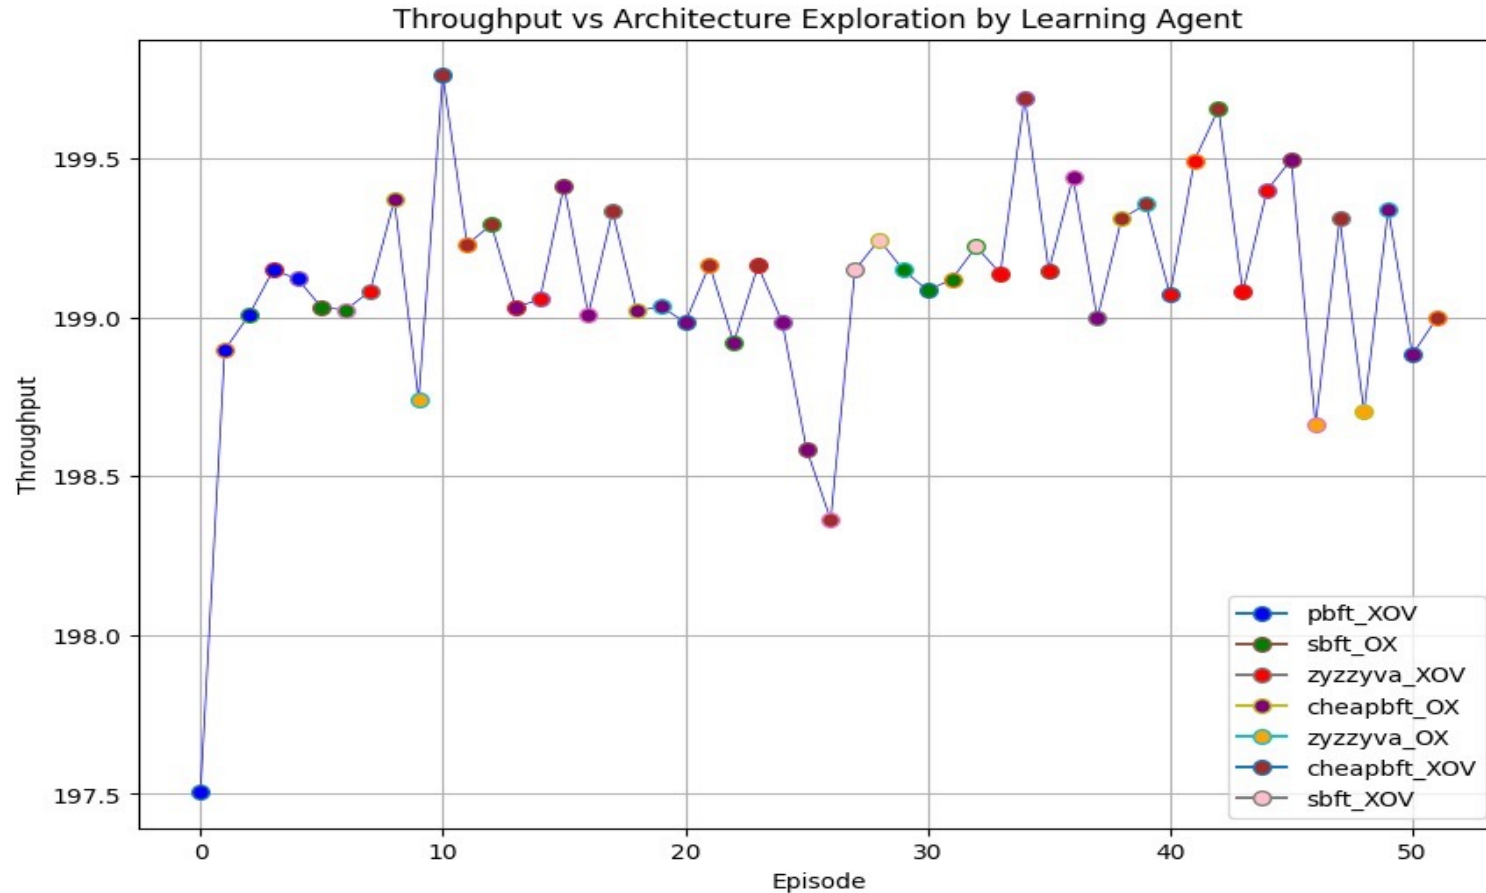
- Our implementation is in a layer on top of the **Bedrock** framework.
- Communication is done over RPC (Client, Learning Agent, Peers)
- **Learning agent:** Python, SkLearn

# Experimental Setup

- Currently running on the local server with 4 nodes.
- As mentioned, each server has its learning agent.
- Each node has a dedicated connection to other nodes



# Evaluation Results – Throughput




Throughput vs Episode

$$\text{Throughput} = \frac{\text{Successful Transactions Count}}{\text{Time}}$$

- The Learning Agent explores by selecting architecture combinations and updating its reward.
- Over 50 episodes, the system demonstrates its ability to adapt by choosing different architectures, and protocols, and retraining itself.



# Evaluation Results - Average execution time



- The Learning Agent explores by selecting architecture combinations and updating its reward.
- Over 50 episodes, the system demonstrates its ability to adapt by choosing different architectures, and protocols, and retraining itself.





# Conclusion

- Fully adaptive system. Adapts to workloads.
- Dynamically selects the top-performing configuration
- Higher throughput compared to fixed configurations



# References

- Adachain - <https://arxiv.org/abs/2211.01580>
- BedRock - <https://www3.cs.stonybrook.edu/~amiri/papers/bedrock.pdf>
- Code : <https://github.com/madhulakkoju/BFTBrain>





# Thank you